

# Supporting architecture evaluation for ATAM scenarios with LLMs

Rafael Capilla<sup>1</sup>, J. Andrés Díaz-Pace<sup>2</sup>, Yamid Ramírez<sup>1</sup>, Jennifer Pérez<sup>3</sup>,  
and Vanessa Rodríguez-Horcajo<sup>3</sup>

<sup>1</sup> Rey Juan Carlos University, Madrid, Spain

`rafael.capilla@urjc.es, ye.ramirez.2024@alumnos.urjc.es`

<sup>2</sup> ISISTAN, CONICET/UNICEN, Tandil, Buenos Aires, Argentina

`andres.diazpace@isistan.unicen.edu.ar`

<sup>3</sup> Universidad Politécnica de Madrid, Spain

`jenifer.perez@upm.es, vanessa.rodriguez.horcajo@upm.es`

**Abstract.** Architecture evaluation methods have long been used to evaluate software designs. Several evaluation methods have been proposed and used to analyze tradeoffs between different quality attributes. Having competing qualities leads to conflicts for selecting which quality-attribute scenarios are the most suitable ones that an architecture should tackle and for prioritizing the scenarios required by the stakeholders. In this context, architecture evaluation is carried out manually, often involving long brainstorming sessions to decide which are the most adequate quality scenarios. To reduce this effort and make the assessment and selection of scenarios more efficient, we suggest the usage of LLMs to partially automate evaluation activities. As a first step to validate this hypothesis, this work studies *MS Copilot* as an LLM tool to analyze quality scenarios suggested by students in a software architecture course and compares the students' results with the assessment provided by the LLM. Our initial study reveals that the LLM produces in most cases better and more accurate results regarding the risks, sensitivity points and tradeoff analysis of the quality scenarios. Overall, the use of generative AI has the potential to partially automate and support the architecture evaluation tasks, improving the human decision-making process.

**Keywords:** Architecture evaluation · Quality-attribute scenarios · Architecture tradeoffs · Large Language Models · Design assistance

## 1 Introduction

For more than two decades, software architects have used different architecture evaluation methods (e.g. ALMA, ATAM ARID, SAAM) [3] to assess quality attributes and determine which and how certain quality-attribute properties should be explicitly considered in the architecture before implementing the system. Addressing quality attributes requires one to identify risks and non-risks of quality scenarios and also the impact of sensitivity points on the architecture [2]. Furthermore, tradeoffs between competing qualities may complicate the selection

and prioritization of scenarios. This activity often leads to long brainstorming sessions between stakeholders and architects.

In order to partially automate the evaluation quality scenarios and double-check if the selected scenarios in an architecture evaluation are adequate, we propose the usage of generative AI techniques like Large Language Models (LLMs). In addition to making brainstorming sessions more effective, our goal is to assist architects in their evaluation tasks by suggesting the most suitable scenarios for improving an architecture based on their pros and cons, alerting architects about risks, and also possible quality-attribute tradeoffs.

In this initial research work, we investigate how an LLM like **MS Copilot** can support quality-driven architecture evaluations. To this end, we adopted a Retrieval Augmented Generation (RAG) strategy [7] to feed an LLM using an architecture assignment from an undergraduate course of computer science students that applied the Architecture Tradeoff Analysis Method (ATAM) [3] and then we use prompts to instruct the LLM to assess the architecture in a stepwise manner.

The remainder of this paper is as follows. Section 2 describes the background to understand the relevant concepts. Section 3 outlines the study design and research goals. Section 4 describes the initial evaluation performed, and Section 5 discusses the conclusions and future work.

## 2 Background

ATAM [6] is a manual method that aims to assess the consequences and impact of architectural decisions motivated by quality-attribute requirements. An ATAM team is responsible for analyzing the quality needs of stakeholders and producing the so-called *utility tree*, which is formed by a set of quality scenarios addressing specific quality concerns. These qualities can impact different system's areas and on the software architecture as well. In addition, each quality scenario is characterized by *stimuli* to which an architecture (or system) must respond and a way to measure the quality-attribute response triggered by the stimuli. The ATAM team must suggest, evaluate, and select the most suitable scenarios by evaluating possible risks, sensitivity points and tradeoffs between qualities.

In particular, the tradeoffs for each scenario reflect the tensions between two or more qualities, thus a criterion must be used to prioritize and select between several scenarios during the brainstorming sessions. An initial experience analyzing tradeoff and sensitivity points using the Analytical Hierarchical Process (AHP) method is described in [5], where the authors rely on pairwise comparisons to support decision-making of commercial off the shelf (COTS) components. The use of AHP to evaluate the quality attributes of the system architecture according to the ISO 25010:2011 quality model<sup>4</sup> is also highlighted in [4]. Only a few approaches, such as those mentioned in [8], support some kind of automation with respect to architecture quality tradeoffs in multi-attribute decision-making.

Using generative AI and LLMs to make decisions is a timely alternative to speed up architects' decisions and refine system architectures [9]. Not far ago,

<sup>4</sup> Now replaced by the ISO 25002:2024 <https://www.iso.org/standard/78175.html>

LLMs have been trained and fine-tuned using datasets [1] in order to respond to natural language questions. Therefore, the use of LLM-based tools with adequate prompts [11] is an important challenge, so that architects do not overlook relevant knowledge when evaluating quality scenarios.

### 3 Study Design

This work evaluates whether generative AI can enhance human decision-making and reflective practices [10] based on prior results of novice architects. We analyzed the results of 9 student projects from a software architecture course at the university level. In this project, a team of students acting as an ATAM team plus another one acting as a customer captured, evaluated, and selected quality scenarios to incorporate certain quality properties in an existing software architecture. Note that although the case study to produce the architecture in a previous assignment was the same, the final architecture produced by each team might exhibit differences. In addition, the quality needs stated by each student acting as the “customer” in the ATAM process can be different in each of the assignments.

The steps that the students performed to apply ATAM were as follows: (i) they used a software architecture designed in a previous assignment as input, and one student acting as the customer indicated a set of quality needs to improve the current architecture, (ii) the quality needs were discussed and refined with 5 students acting as the ATAM team, (iii) the students followed the ATAM method to create a utility tree describing quality scenarios for each of the quality attributes derived from the system quality goals and affecting different parts of the existing architecture, (iv) each scenario was evaluated to identify possible risks, sensitivity points, and tradeoffs between quality attributes affecting the scenario and (v) the final scenarios selected by students were engineered in the architecture. All the information from the students’ assignments was gathered in PDF documents. We provided these documents as input to **MS Copilot**<sup>5</sup> and used a RAG strategy<sup>6</sup> [7] to evaluate the degree to which the scenarios chosen by the students were adequate. An example of the results generated by the LLM using the sequence of four prompts is given here<sup>7</sup>.

The co-authors tested the validity and accuracy of the scenarios using four different prompts for **MS Copilot**. The steps of the ATAM process to evaluate and select the quality scenarios are depicted in Figure 1, enriched with the 4 prompts used for **MS Copilot**. In particular, we prompted the LLM to identify possible risks, sensitivity points and perform a tradeoff analysis. Afterwards, we asked **MS Copilot** to perform a selection of scenarios based on the previous information. From the 9 student groups evaluated, we discarded two groups

<sup>5</sup> <https://copilot.microsoft.com/>

<sup>6</sup> RAG integrates external domain-specific sources into the LLM to improve the accuracy and relevance of the response

<sup>7</sup> <https://anonymous.4open.science/r/archevaluation-llms-1FDE/prompting-example.md>

because the quality of the scenarios was rather low and they provided several unclear scenarios.

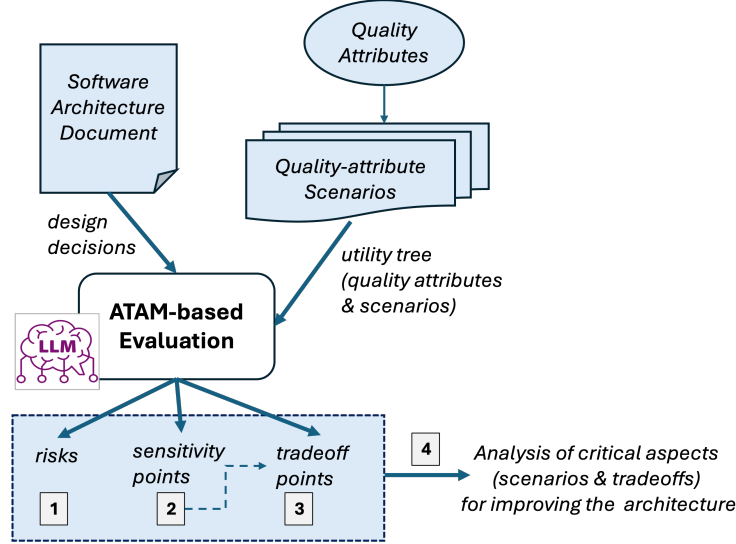


Fig. 1: Inputs and outputs of the ATAM process, including the steps where the LLM prompts were used (1, 2, 3, 4)

## 4 Preliminary Evaluation

In this section we describe the results of our LLM-approach to simulate the ATAM process. The scenarios selected by the students were based on their analysis of risks for the scenarios, identification of sensitivity points in the architecture that could affect the scenarios, and a tradeoff analysis between quality attributes. Thus, we used four prompts to query *MS Copilot* about these aspects to validate the scenarios, in addition to asking the LLM to suggest scenarios. We provide a reproducibility kit that contains the input artifacts and prompts.

### Prompts #1 and #2. Identification of risks and sensitivity points:

In this step we asked *MS Copilot* to infer risks related to all the scenarios described by the students. In the context of ATAM, we refer to a risk as a problematic architectural decision that can potentially lead to negative consequences if not addressed properly. A sensitivity point, in turn, is a design decision that affects the achievement of a particular quality-attribute scenario.

The results are shown in Table 1, which distinguishes the risks and sensitivity points found by the students and by *MS Copilot*. As we can observe, *MS Copilot* was extremely useful in detecting additional risks and sensitivity points (i.e., items not identified by the students), which are relevant for the selection of scenarios. In several groups, there was a good intersection (e.g. *G6*) while in other groups the LLM detected more risks and sensitivity points (e.g., *G1*, *G3*, *G8*). In those cases where *MS Copilot* suggested additional risks, we list such

scenarios in the last column of Table 1 in order to indicate that in some cases certain scenarios should not have been chosen by the students. However, this decision depends on factors such as the risk severity, probability of occurrence, or business importance of the scenario, because not all risks are equally critical. If a risk identified for a chosen scenario is not critical, we assume that selecting such scenario would not harm the system. In such situations, we could ask MS Copilot to prioritize the risks to perform a more accurate scenario selection.

Table 1: Risks and sensitivity points

Groups	Risks (students)	Risks (Copilot)	Sensitivity points (students)	Sensitivity points (Copilot)	Scenarios selected with risks
G1	1.1, 1.4, 3.1, 4.1	1.1, 1.4, 2.2, 2.3, 3.1, 3.3, 4.1, 4.4, 6.3	Database, backup security, database scalability	Database, backup security, database scalability, Authentication, data cache, system workload, horizontal scalability, network security	3.3
G2	1.1, 1.3, 1.4, 2.1, 2.4, 3.1, 3.3, 3.4, 3.5	1.1, 1.2, 1.4, 1.5, 2.1, 2.3, 2.4, 3.1, 3.3, 3.4	N/A	Data security, database performance, server concurrency, database replication, hot updates	1.2, 2.3, 3.3
G3	1.1, 1.2, 1.3, 1.4, 1.5, 2.1, 2.2, 2.3, 3.1, 3.2, 3.3, 3.4, 4.1, 4.2, 4.3	1.1, 1.2, 1.3, 1.4, 1.5, 2.1, 2.2, 2.3, 3.1, 3.2, 3.3, 3.4, 4.1, 4.2, 4.3	Peak workload, latency, bottlenecks, database security, biometric security	Peak workload, latency, bottlenecks, database security, biometric security, network capacity, complexity of authentication methods, lack of infrastructure redundancy, integration between components	1.1, 2.3, 3.2
G4	1.1, 1.6, 2.1, 2.4	1.1, 1.2, 1.6, 2.1, 2.4, 3.1	Database	Database, encryption mechanism, load balancer, cache, 2FA multi-factor authentication	1.1, 1.6
G6	1.1, 2.2, 2.4, 3.2	1.1, 2.2, 2.4, 3.2	Database	Database	None
G7	1.1, 1.2, 3.1, 3.2	1.1, 1.2, 2.3, 3.1, 3.2	Database encryption, database scalability	Database encryption, database scalability	3.2
G8	1.1, 1.2, 1.3, 2.2, 2.3, 3.1, 3.6	1.1, 1.2, 1.3, 2.2, 2.3, 3.1, 3.4, 3.5, 3.6	Server availability, database caching, database encryption	Database caching, database encryption, authentication failures, server monitoring	1.3, 2.2, 3.1, 3.4

**Prompt #3. Tradeoff analysis:** In this step we prompted MS Copilot to perform a quality-attribute tradeoff analysis for the scenarios identified in the utility tree by each group. The results are shown in Table 2. The first two columns indicate the qualities involved in each tradeoff identified by the students and by MS Copilot, while the third and fourth columns indicate the scenarios involved and selected during the tradeoff analysis by the students and MS Copilot as well. As it can be seen, MS Copilot identified more trade-offs between qualities in all cases. Therefore, the analysis performed by the LLM involves more scenarios for the tradeoff analysis. Also, in general terms, the scenarios chosen by MS Copilot are in line with the selection made by the students except in some cases, where MS Copilot considered that more scenarios could be selected (e.g. *G6*, *G7*).

Table 2: Tradeoff analysis points

Groups	Tradeoff analysis (students)	Tradeoff analysis (Copilot)	Scenarios affected by tradeoffs	Winning scenarios after tradeoff
G1	Performance-Security, Performance-Reliability	Performance-Security, Performance-Reliability, Security-Usability, Performance-Scalability, Security-Maintainability	Students: None Copilot: 1.1, 1.2, 2.2, 2.3, 2.4, 4.1, 4.2, 5.2	Students: 1.2, 2.1, 3.3, 4.2, 5.2, 6.6 Copilot: 1.2, 2.1, 3.3, 4.2, 5.2
G2	Security-Performance, Security-Availability	Security-Performance, Security-Availability, Performance-Availability	Students: 1.1, 1.3, 1.6 Copilot: 1.1, 1.3, 2.1, 2.3, 2.4, 2.6, 3.1, 3.6	Students: 1.2, 1.6, 2.3, 2.6, 2.9, 3.3, 3.6 Copilot: 1.2, 1.6, 2.3, 2.6, 2.9, 3.6
G3	Scalability-Performance, Security-Performance, Availability-Performance, Authenticity-Compatibility	Scalability-Latency, Security-Performance, Availability-Complexity, Authenticity-Usability	Students: ALL Copilot: ALL	Students: 1.1, 2.3, 3.2, 4.4 Copilot: 1.2, 2.3, 3.1, 4.3
G4	Security-Performance, Scalability-Performance	Security-Performance, Scalability-Performance, Scalability-Complexity, Performance-Consistency, Security-Usability, Availability-Complexity	Students: 1.1, 1.6, 2.2, 2.3, 3.1 Copilot: 1.2, 2.1, 2.4, 2.5, 3.1	Students: 1.1, 1.6, 2.2, 2.5, 3.1 Copilot: 1.2, 2.1, 2.4, 3.1
G6	Performance-Interoperability, Security-Performance, Performance-Availability	Security-Performance, Interoperability-Complexity, Availability-Cost, Scalability-Simplicity	Students: ALL Copilot: ALL	Students: 1.2, 2.1, 3.4 Copilot: 1.2, 2.1, 2.3, 3.4
G7	Security-Performance	Security-Performance, Performance-Scalability, Scalability-Cost, Security-Usability, Performance-Complexity, Scalability-Maintainability	Students: 1.1, 1.2, 1.3 Copilot: ALL	Students: 1.3 Copilot: 1.2, 1.3, 2.3, 3.1, 3.3
G8	Reliability-Availability, Reliability-Performance Efficiency, Security-Performance	Reliability-Performance Efficiency, Security-Performance Efficiency, Scalability-Maintainability, Availability-Cost, Security-Usability	Students: ALL Copilot: 1.1, 1.2, 1.3, 2.1, 2.2, 2.3, 3.1, 3.2, 3.3, 3.4	Students: 1.3, 2.1, 2.2, 2.4 Copilot: 1.2, 1.3, 2.1, 3.1

**Prompt #4. Selection of scenarios:** In Table 3 we show, for each group and quality attribute, the identifiers of the scenarios selected by the students and by MS Copilot. We marked in red those cases where the scenarios chosen by the students and by the LLM were very different, and in green those that exhibited a very good match. In the rest of the cases, there was some degree of scenario matching, but also some differences between the students and MS Copilot. We must remark that in the case of G6 the matching of the scenarios was almost perfect. Overall, only in 4 cases out of 30 the scenarios selected by MS Copilot diverged from the selection of the students. Thus, in general terms, we can consider that the use of an LLM was beneficial to evaluate and select scenarios. Finally, we asked MS Copilot to select the most suitable scenarios for a specific quality attribute and area of the system affected by it, but the LLM often selected only one scenario. This is an example of the challenges that still exist when asking an LLM to evaluate and select quality scenarios.

In the case of the risks and sensitivity points identified by the students and MS Copilot in Table 1, MS Copilot was able to identify additional risks and sensitivity points in most cases affecting more scenarios than those of the stu-

Table 3: Quality-attribute scenarios selected

Groups and scenarios	Quality 1	Quality 2	Quality 3	Quality 4	Quality 5	Quality 6	Quality 7
G1	Security	Security	Reliability	Performance Efficiency (data access)	Performance Efficiency (orders)		
Sc. students	1.2	1.2	3.3, 4.2	5.2	6.2		
Sc. Copilot	1.2	2.1	3.3, 4.2	5.2	None		
G2	Security (database)	Security (server)	Performance (database)	Performance (server)	Performance (orders)	Availability (database)	Availability (orders)
Sc. students	1.2	1.6	2.3	2.6	2.9	3.3	3.6
Sc. Copilot	1.2, 1.3	1.5, 1.6	2.2, 2.3	2.5	2.8, 2.9	3.1, 3.2	3.4, 3.6
G3	Scalability	Security (database)	Security-Authenticity	Availability			
Sc. students	1.2	2.3	3.2	4.1, 4.2			
Sc. Copilot	1.1, 1.2	2.2, 2.3	3.1, 3.2	4.3, 4.4			
G4	Security (database)	Security (client)	Scalability (GW)	Scalability (database)	Performance (database)		
Sc. students	1.1j	1.6	2.2	2.5	3.1		
Sc. Copilot	1.1	1.4	2.2	2.5	3.1		
G6	Interoperability	Security	Availability				
Sc. students	1.2	2.1	3.4				
Sc. Copilot	1.2	2.1	3.4				
G7	Security	Performance	Scalability				
Sc. students	1.3	2.2, 2.3	3.1				
Sc. Copilot	1.2, 1.3	2.2, 2.3	3.2, 3.3				
G8	Reliability	Performance Efficiency	Security				
Sc. students	1.3	2.1, 2.2	3.1, 3.2, 3.4				
Sc. Copilot	1.3	2.1	3.2, 3.3				

dents. However, not all the new risks identified by **MS Copilot** had the same degree of severity. Some risks are important, while others can be overlooked in the selection of a given scenario. The last column of the table indicates scenarios selected by the students but affected by risks identified by the **MS Copilot**. Due to space constraints, we cannot provide a detailed classification of the severity of the risks for each scenario. Therefore, a thorough analysis of the severity of each risk must be made in a further study. In addition, **MS Copilot** identified new sensitivity points in the architecture that the students did not raise. This is particularly important for identifying potential quality tradeoffs, hidden risks, or extending the students' analysis by including critical system design areas.

With regard to the tradeoff analysis results shown in Table 2, **MS Copilot** identified in all cases more tradeoffs than the students. Nevertheless, some of these tradeoffs belong to certain qualities (e.g., cost, usability) that the students were told to ignore due to the lack of cost numbers or the difficulty to evaluate certain qualities (e.g., usability). Although **MS Copilot** was able to suggest more scenarios affected by these tradeoffs, a deeper analysis would be recommended to assess the accuracy of the tradeoffs. Finally, from the scenarios selected from the tradeoff analysis, our initial results reveal that **MS Copilot** matched the students' results in several cases, while in others suggested different scenarios. For example, in group *G6*, the students reported only one scenario (1.3) for a tradeoff between two qualities while **MS Copilot** indicated 5 scenarios from tradeoffs between 6 pairs of qualities. The results of the students are independent

of the final numbers of scenarios selected. In other cases, such as group *G3*, all scenarios except one are different in the tradeoff analysis between the students and *MS Copilot*, probably because the LLM identified 6 tradeoffs versus the two tradeoffs reported by the students.

## 5 Conclusions and Future Work

In this paper, we have investigated how LLMs can provide support and improve the results of an ATAM architecture evaluation process carried out by students of a software architecture course. Overall, *MS Copilot* was able to produce more information for key phases of ATAM. More specifically, the LLM produced more results in the identification of risks, sensitivity points and tradeoff analysis, which contributes to improve the outputs of the evaluation process.

The results so far confirm our hypothesis about the usefulness of LLMs in the assessment of ATAM scenarios, identifying additional information to support the architecture evaluation of quality requirements. In future work, we plan to explore two paths of work: (i) using LLMs to create and rank quality scenarios, (ii) performing a deeper analysis of the role of concrete risks, sensitivity points, and accuracy of tradeoffs. Finally, alternative LLMs can be explored to cross-validate the results of *MS Copilot* and to see whether contextual knowledge from previous evaluations can affect the reported results.

**Data Availability** We provide an anonymized site with the data used in the paper: <https://anonymous.4open.science/r/archevaluation-llms-1FDE/>.

## Bibliography

- [1] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J.D., Dhariwal, P., Nee-lakantan, A., Shyam, P., Sastry, G., Askell, A., et al.: Language models are few-shot learners. *Adv Neural Inf Process Sys* 33, 1877–1901 (2020)
- [2] Cervantes, H., Kazman, R.: *Designing Software Architectures: A Practical Approach* 2nd Edition. SEI series in software engineering, Addison-Wesley (2024)
- [3] Clements, P., Kazman: *Evaluating Software Architectures: Methods and Case Studies*. SEI series in software engineering, Addison-Wesley (2001)
- [4] Darwish, M., Shehab, E.: Framework for engineering design systems architectures evaluation and selection: Case study. *Procedia CIRP* 60, 128–132 (2017), complex Systems Engineering and Development Proceedings of the 27th CIRP Design Conference Cranfield University, UK 10th – 12th May 2017
- [5] Ibrahim, H., Far, B.H., Eberlein, A.: Tradeoff and sensitivity analysis of a hybrid model for ranking commercial off-the-shelf products. In: 2009 16th Annual IEEE International Conference and Workshop on the Engineering of Computer Based Systems. pp. 119–127 (2009)
- [6] Kazman, R., Klein, M., Clements, P.: *Atam:sm method for architecture evaluation*. Tech. Rep. CMU/SEI-2000-TR-004, Carnegie Mellon University - Software Engineering Institute (2000)
- [7] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.t., Rocktäschel, T., et al.: Retrieval-augmented generation for knowledge-intensive nlp tasks. *Adv Neural Inf Process Syst* 33, 9459–9474 (2020)



- [8] Lytra, I., Carrillo, C., Capilla, R., Zdun, U.: Quality attributes use in architecture design decision methods: research and practice. *Computing* 102(2), 551–572 (2020)
- [9] Ozkaya, I.: Can architecture knowledge guide software development with generative ai? *IEEE Software* 40(5), 4–8 (2023)
- [10] Razavian, M., Tang, A., Capilla, R., Lago, P.: In two minds: how reflections influence software design thinking. *J. Softw.: Evol. Process* 28(6), 394–426 (2016)
- [11] White, J., Hays, S., Fu, Q., Spencer-Smith, J., Schmidt, D.C.: Chatgpt prompt patterns for improving code quality, refactoring, requirements elicitation, and software design. *CoRR* abs/2303.07839 (2023)